

Karthik Pradeep Hegadi

2KE20CS032

Assignment 18

Understood. To follow the provided instructions and create the files/directory using the same name and case as provided in the task steps, please provide me with the specific names and case instructions for the files/directory you want to create.

STORAGE

Assignment 1- :Test Disk Performance OPS With Fio utility for random read write Operations.

Instructions

Run this command :

(fio—randrepeat=1—ioengine=libaio—direct=1--gtodreduce=1--name=test--filename=test)

Or

(fio --randrepeat=1 --ioengine=libaio --direct=1 --gtod_reduce=1 --name=/path/to/testfile -- filename=/path/to/testfile --bs=2k --iodepth=64 --size=2G --readwrite=randrw --rwmixread=75)

This command will create a 2 GB file(-size option), and perform 2KB reads and writes(--bs option) using 75%/25%, with 64 operations(-iodepth option) running at a time. You can change these value and test as well.

To measure random write OPS use option:

--readwrite=randwrite

To measure random read OPS use option:

--readwrite=randread

Check & record the output of this command to know the disk performance in terms of read operations per second (IOPS) and write operations per second.

What is fio?

fio is a tool that will spawn a number of threads or processes doing a particular type of I/O action as specified by the user. The typical use of fio is to write a job file matching the I/O load one wants to simulate

Parameters:

iodepth: Number of I/O units to keep in flight against the file.

bs: filesystem block size

randrepeat: Seed the random number generator used for random I/O patterns in a predictable way so the pattern is repeatable across runs. **Default:** true. Value 1 is false

ioengine: Defines how the job issues I/O to the file. libaio is Linux native asynchronous I/O

-direct: direct or buffered

-god_reduce: tod(time of day). Whether to reduce gettimeofday call during the operation

-filename: name of the file created. You can see in current directory

--bs: Block Size

-odepth: Number of I/O units to keep in flight against the file.

-size: Size of the file to be created for test

-readwrite: Type of I/O pattern. "randrw" Random mixed reads and writes.

-rwmixread: Percentage of a mixed workload that should be reads. Default: 50.

```
~ (0.239s)
ssh lusy@10.211.55.6

lusy@localhost.localdomain ~ (3.226s)
sudo yum install fio
[sudo] password for lusy:
Last metadata expiration check: 1:17:57 ago on Sat 19 Aug 2023 10:41:06 AM IST.
Package fio-3.35-1.el9.aarch64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
```

```

lusi@localhost.localdomain ~ (18.835s)
sudo fio --randrepeat=1 --ioengine=libaio --direct=1 --gtod_reduce=1 --name=/path/to/testfile --filename=/path/to/testfile --bs=2k
-rwmixread=75
/path/to/testfile: (g=0): rw=randrw, bs=(R) 2048B-2048B, (W) 2048B-2048B, (T) 2048B-2048B, ioengine=libaio, iodepth=64
fio-3.35
Starting 1 process
/path/to/testfile: Laying out IO file (1 file / 2048MiB)
Jobs: 1 (f=1): [m(1)][100.0%][r=84.1MiB/s,w=28.0MiB/s][r=43.1k,w=14.3k IOPS][eta 00m:00s]
/path/to/testfile: (groupid=0, jobs=1): err= 0: pid=26680: Sat Aug 19 12:31:50 2023
  read: IOPS=47.2k, BW=92.2MiB/s (96.7MB/s)(1535MiB/16645msec)
    bw ( Kib/s): min=14416, max=157708, per=100.0%, avg=94498.24, stdev=45405.00, samples=33
    iops : min= 7208, max=78854, avg=47249.09, stdev=22702.51, samples=33
  write: IOPS=15.8k, BW=30.8MiB/s (32.3MB/s)(513MiB/16645msec); 0 zone resets
    bw ( Kib/s): min= 4744, max=52852, per=100.0%, avg=31576.09, stdev=15202.05, samples=33
    iops : min= 2372, max=26426, avg=15788.00, stdev=7601.05, samples=33
  cpu : usr=5.49%, sys=33.34%, ctx=74890, majf=0, minf=17
  IO depths : 1=0.1%, 2=0.1%, 4=0.1%, 8=0.1%, 16=0.1%, 32=0.1%, >=64=100.0%
    submit : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.0%, >=64=0.0%
    complete : 0=0.0%, 4=100.0%, 8=0.0%, 16=0.0%, 32=0.0%, 64=0.1%, >=64=0.0%
  issued rwt: total=785920,262656,0,0 short=0,0,0 dropped=0,0,0
  latency : target=0, window=0, percentile=100.00%, depth=64

Run status group 0 (all jobs):
  READ: bw=92.2MiB/s (96.7MB/s), 92.2MiB/s-92.2MiB/s (96.7MB/s-96.7MB/s), io=1535MiB (1610MB), run=16645-16645msec
  WRITE: bw=30.8MiB/s (32.3MB/s), 30.8MiB/s-30.8MiB/s (32.3MB/s-32.3MB/s), io=513MiB (538MB), run=16645-16645msec

Disk stats (read/write):
  dm-0: ios=784151/262062, merge=0/0, ticks=596440/266730, in_queue=863170, util=99.53%, aggrios=785109/262600, aggrmerge=963/10
  946, aggrutil=99.02%
  sda: ios=785109/262600, merge=963/100, ticks=598438/263453, in_queue=861946, util=99.02%

```

Assignment 2- Setup NFS server in one linux system and access the shared directory from another linux system which will act as NFS client.

Refer Network file system slide for this Assignment

Note- Here to comlete this assignment I have cloned two machines

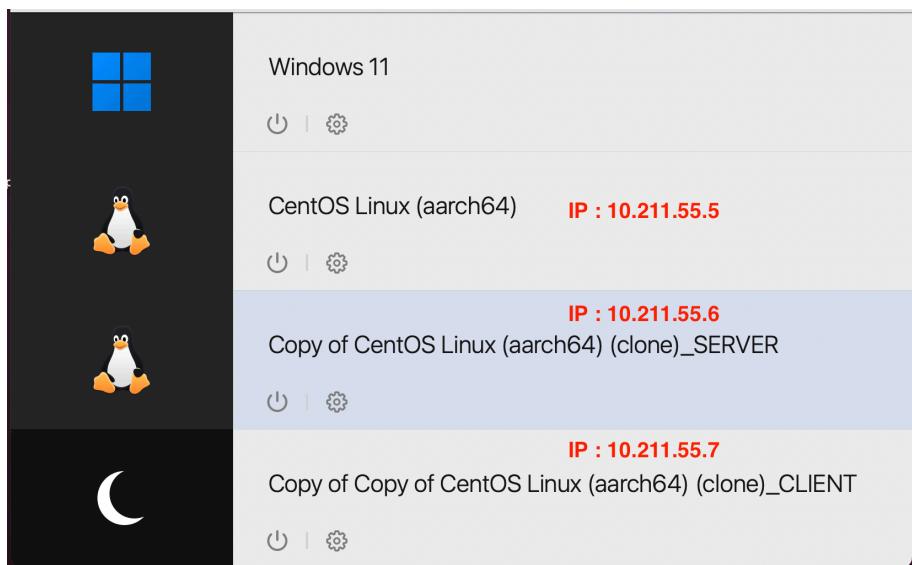
image 1 : which is orignale image of centos.

image 2 : which i took from origmal machine :**Clone 1**

image 3 : which I took form clone 1 machine :**Clone 2**

For Clone 1: Machine I have added new user called **SERVER** (also added sudo previledge to it).

For Clone 2: Machine I have added new user called **CLIENT** (also added sudo previledge to it).



SERVER configuration:

```
lusy@localhost.localdomain ~ (18m 44.49s)
su SERVER
Password:
[SERVER@localhost lusy]$ cd
[SERVER@localhost ~]$ sudo yum install nfs-utils
[sudo] password for SERVER:
Last metadata expiration check: 3:32:45 ago on Sat 19 Aug 2023 10:41:06 AM IST.
Package nfs-utils-1:2.5.4-20.el9.aarch64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
```

```
[SERVER@localhost ~]$ sudo systemctl enable nfs-server
Created symlink /etc/systemd/system/multi-user.target.wants/nfs-server.service → /usr/lib/systemd/system/nfs-server.service
```

```
[SERVER@localhost ~]$ sudo systemctl start nfs-server
[SERVER@localhost ~]$ sudo systemctl status nfs-server
● nfs-server.service - NFS server and services
  Loaded: loaded (/usr/lib/systemd/system/nfs-server.service; enabled; preset: enabled)
  Drop-In: /run/systemd/generator/nfs-server.service.d
            └─order-with-mounts.conf
    Active: active (exited) since Sat 2023-08-19 13:30:51 IST; 43min ago
      Main PID: 18690 (code=exited, status=0/SUCCESS)
        CPU: 13ms

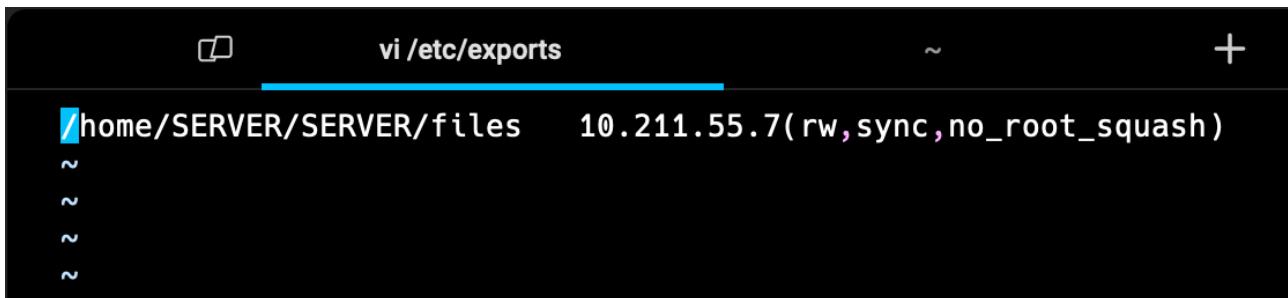
Aug 19 13:30:51 localhost.localdomain systemd[1]: Starting NFS server and services.
Aug 19 13:30:51 localhost.localdomain systemd[1]: Finished NFS server and services.
```

```
[SERVER@localhost ~]$ sudo systemctl enable rpcbind
[SERVER@localhost ~]$ sudo systemctl start rpcbind
[SERVER@localhost ~]$ sudo systemctl status rpcbind
● rpcbind.service - RPC Bind
  Loaded: loaded (/usr/lib/systemd/system/rpcbind.service; enabled; preset: enabled)
  Active: active (running) since Sat 2023-08-19 13:30:51 IST; 43min ago
  TriggeredBy: ● rpcbind.socket
    Docs: man:rpcbind(8)
  Main PID: 18657 (rpcbind)
    Tasks: 1 (limit: 10677)
   Memory: 1.2M
     CPU: 31ms
    CGroup: /system.slice/rpcbind.service
              └─18657 /usr/bin/rpcbind -w -f

Aug 19 13:30:51 localhost.localdomain systemd[1]: Starting RPC Bind...
Aug 19 13:30:51 localhost.localdomain systemd[1]: Started RPC Bind.
```

```
[SERVER@localhost ~]$ cd
[SERVER@localhost ~]$ ls
SERVER
[SERVER@localhost ~]$ cd SERVER
[SERVER@localhost SERVER]$ ls
files
[SERVER@localhost SERVER]$ cd files
[SERVER@localhost files]$ pwd
/home/SERVER/SERVER/files
```

```
[SERVER@localhost files]$ sudo vi /etc/exports
```



The screenshot shows a terminal window with the title "vi /etc/exports". The file contains one line of text: "/home/SERVER/SERVER/files 10.211.55.7(rw,sync,no_root_squash)". There are four blank lines below this entry.

```
/home/SERVER/SERVER/files 10.211.55.7(rw,sync,no_root_squash)
```

```
[SERVER@localhost files]$ sudo systemctl stop firewalld.service
[SERVER@localhost files]$ sudo systemctl status firewalld.service
● firewalld.service - firewalld - dynamic firewall daemon
   Loaded: loaded (/usr/lib/systemd/system/firewalld.service; enabled; preset: en>
   Active: inactive (dead) since Sat 2023-08-19 14:18:34 IST; 11s ago
     Duration: 1h 3min 53.949s
       Docs: man:firewalld(1)
    Process: 691 ExecStart=/usr/sbin/firewalld --nofork --nopid $FIREWALLD_ARGS (co>
 Main PID: 691 (code=exited, status=0/SUCCESS)
      CPU: 279ms

Aug 19 13:14:40 localhost systemd[1]: Starting firewalld - dynamic firewall daemon.>
Aug 19 13:14:40 localhost systemd[1]: Started firewalld - dynamic firewall daemon.
Aug 19 14:18:34 localhost.localdomain systemd[1]: Stopping firewalld - dynamic fire>
Aug 19 14:18:34 localhost.localdomain systemd[1]: firewalld.service: Deactivated su>
Aug 19 14:18:34 localhost.localdomain systemd[1]: Stopped firewalld - dynamic firew>
```

```
[SERVER@localhost files]$ ip a
1: lo: <LOOPBACK,UP,LOWER_UP> mtu 65536 qdisc noqueue state UNKNOWN group default
  link/loopback 00:00:00:00:00:00 brd 00:00:00:00:00:00
  inet 127.0.0.1/8 scope host lo
    valid_lft forever preferred_lft forever
    inet6 ::1/128 scope host
      valid_lft forever preferred_lft forever
2: enp0s5: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc fq_codel state UP group default
  link/ether 00:1c:42:cd:4f:6f brd ff:ff:ff:ff:ff:ff
  inet 10.211.55.6/24 brd 10.211.55.255 scope global dynamic noprefixroute enp0s5
    valid_lft 1462sec preferred_lft 1462sec
    inet6 fdb2:2c26:f4e4:0:21c:42ff:fe:fed:4f6f/64 scope global dynamic noprefixroute
      valid_lft 2591911sec preferred_lft 604711sec
    inet6 fe80::21c:42ff:fe:fed:4f6f/64 scope link noprefixroute
      valid_lft forever preferred_lft forever
```

CLIENT configuration:

```
lusy@localhost.localdomain ~ (13m 1.23s)
su CLIENT
Password:
[CLIENT@localhost lusy]$ cd
[CLIENT@localhost ~]$ sudo yum install nfs-utils
[sudo] password for CLIENT:
Last metadata expiration check: 3:38:30 ago on Sat 19 Aug 2023 10:41:06 AM IST.
Package nfs-utils-1:2.5.4-20.el9.aarch64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
```

```
Complete!
[CLIENT@localhost ~]$ ls
mounted
[CLIENT@localhost ~]$ cd mounted
[CLIENT@localhost mounted]$ ls
mountedfiles
[CLIENT@localhost mounted]$ cd mountedfiles
[CLIENT@localhost mountedfiles]$ pwd
/home/CLIENT/mounted/mountedfiles
```

```
[CLIENT@localhost ~]$ sudo mount 10.211.55.6:/home/SERVER/SERVER/files /home/CLIENT/mounted/mountedfiles
```

```
[CLIENT@localhost ~]$ ls
mounted
[CLIENT@localhost ~]$ cd mounted
[CLIENT@localhost mounted]$ ls
mountedfiles
[CLIENT@localhost mounted]$ cd mountedfiles
[CLIENT@localhost mountedfiles]$ ls
1
[CLIENT@localhost mountedfiles]$ touch files1.txt
```

```
[CLIENT@localhost mountedfiles]$ vi files1.txt
[CLIENT@localhost mountedfiles]$ ls
files1.txt
[CLIENT@localhost mountedfiles]$ cat files1.txt
WELCOME TO NFS
[CLIENT@localhost mountedfiles]$ █
```

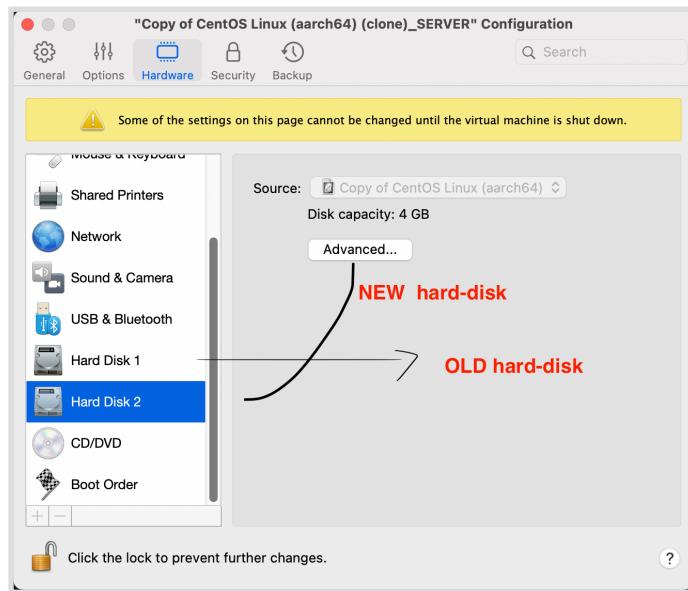
Verification:

```
[SERVER@localhost SERVER]$ cd files
[SERVER@localhost files]$ ls
1 files1.txt
[SERVER@localhost files]$ cat files1.txt
WELCOME TO NFS
[SERVER@localhost files]$
```

Assignment 3 - Add new Virtual hard disk Drive for the virtual box and partition the drive in linux then create file system and mount the file system.

Instructions

1. Add new hard-disk I am using parallel desktop. I followed these instructions (<https://kb.parallels.com/117649>)



1.Run "lsblk" command to check the partitions

```
lusy@localhost.localdomain ~ (0.03s)
lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    0   64G  0 disk
└─sda1     8:1    0  600M  0 part /boot/efi
└─sda2     8:2    0   1G  0 part /boot
└─sda3     8:3    0 62.4G  0 part
  ├─cs-root 253:0  0 40.6G  0 lvm  /
  ├─cs-swap 253:1  0   2G  0 lvm  [SWAP]
  └─cs-home 253:2  0 19.8G  0 lvm  /home
sdb        8:16   0   4G  0 disk
sr0       11:0   1 1024M  0 rom
```

2.create a new partition using the command “fdisk/dev/sdb”

```
lusy@localhost.localdomain ~ (2m 19.95s)
sudo fdisk /dev/sdb
[sudo] password for lusy:

Welcome to fdisk (util-linux 2.37.4).
Changes will remain in memory only, until you decide to write them.
Be careful before using the write command.

Device does not contain a recognized partition table.
Created a new DOS disklabel with disk identifier 0x2b848c38.

Command (m for help):
```



```
Command (m for help): m
```

Help:

```
DOS (MBR)
a    toggle a bootable flag
b    edit nested BSD disklabel
c    toggle the dos compatibility flag

Generic
d    delete a partition
F    list free unpartitioned space
l    list known partition types
n    add a new partition
p    print the partition table
t    change a partition type
v    verify the partition table
i    print information about a partition

Misc
m    print this menu
u    change display/entry units
x    extra functionality (experts only)

Script
I    load disk layout from sfdisk script file
O    dump disk layout to sfdisk script file

Save & Exit
w    write table to disk and exit
q    quit without saving changes

Create a new label
g    create a new empty GPT partition table
G    create a new empty SGI (IRIX) partition table
o    create a new empty DOS partition table
s    create a new empty Sun partition table
```

3. **p** option will show the list of partitions

```
Command (m for help): p

Disk /dev/sdb: 4 GiB, 4294967296 bytes, 8388608 sectors
Disk model: Copy of CentOS L
Units: sectors of 1 * 512 = 512 bytes
Sector size (logical/physical): 512 bytes / 4096 bytes
I/O size (minimum/optimal): 4096 bytes / 4096 bytes
Disklabel type: dos
Disk identifier: 0x2b848c38
```

4. **n** option will create new partition

5. Create 1 primary partition

6. specify where the partition will begin and end by accepting the default values.

```
Command (m for help): n
Partition type
  p  primary (0 primary, 0 extended, 4 free)
  e  extended (container for logical partitions)
Select (default p): p
Partition number (1-4, default 1): 1
First sector (2048-8388607, default 2048):
Last sector, +/-sectors or +/-size{K,M,G,T,P} (2048-8388607, default 8388607):

Created a new partition 1 of type 'Linux' and of size 4 GiB.

Command (m for help): w
```

7. **w** option will save the changes

(Note: Since this is the first partition we can start at the first available sector and since we want to use the entire disk we can specify the last sector as the end. Note that, if you wish to create multiple partitions - you can even specify the size of each partition by sectors, bytes, kilobytes or megabytes)

```
Command (m for help): w
The partition table has been altered.
Calling ioctl() to re-read partition table.
Syncing disks.
```

Verification:

```
lusy@localhost.localdomain ~ (0.03s)
lsblk
NAME      MAJ:MIN RM  SIZE RO TYPE MOUNTPOINTS
sda        8:0    0   64G  0 disk
└─sda1     8:1    0  600M  0 part /boot/efi
└─sda2     8:2    0   1G  0 part /boot
└─sda3     8:3    0 62.4G  0 part
  ├─cs-root 253:0  0 40.6G  0 lvm  /
  ├─cs-swap 253:1  0   2G  0 lvm  [SWAP]
  └─cs-home 253:2  0 19.8G  0 lvm  /home
sdb        8:16   0    4G  0 disk
└─sdb1     8:17   0    4G  0 part
sr0       11:0   1 1024M  0 rom
```

8.Create the file system on the partition using the command “/usr/sbin/mkfs.ext4 -L data /dev/sdb1”

```
lusy@localhost.localdomain ~ (0.104s)
sudo /usr/sbin/mkfs.ext4 -L data /dev/sdb1

mke2fs 1.46.5 (30-Dec-2021)
Discarding device blocks: done
Creating filesystem with 1048320 4k blocks and 262144 inodes
Filesystem UUID: 61ae6cb2-8cde-49c1-a176-115218a47ef9
Superblock backups stored on blocks:
    32768, 98304, 163840, 229376, 294912, 819200, 884736

Allocating group tables: done
Writing inode tables: done
Creating journal (16384 blocks): done
Writing superblocks and filesystem accounting information: done
```

9. Create the directory /data

```
lusy@localhost.localdomain ~ (0.06s)
sudo mkdir /data
```

10. Mount the file system using the command "mount LABEL=data /data"

```
lusy@localhost.localdomain ~ (0.11s)
sudo mount LABEL=data /data
```

11.Run the mount command with no arguments which shows all currently mounted filesystems.
You can use the below command as well.

```
lusy@localhost.localdomain ~ (0.027s)
mount

proc on /proc type proc (rw,nosuid,nodev,noexec,relatime)
sysfs on /sys type sysfs (rw,nosuid,nodev,noexec,relatime,seclabel)
devtmpfs on /dev type devtmpfs (rw,nosuid,seclabel,size=4096k,nr_inodes=213567,mode=0755)
securityfs on /sys/kernel/security type securityfs (rw,nosuid,nodev,noexec,relatime)
tmpfs on /dev/shm type tmpfs (rw,nosuid,nodev,seclabel,inode64)
devpts on /dev/pts type devpts (rw,nosuid,noexec,relatime,seclabel,gid=5,mode=620)
tmpfs on /run type tmpfs (rw,nosuid,nodev,seclabel,size=349380k,nr_inodes=819200,mode=0755)
cgroup2 on /sys/fs/cgroup type cgroup2 (rw,nosuid,nodev,noexec,relatime,seclabel)
pstore on /sys/fs/pstore type pstore (rw,nosuid,nodev,noexec,relatime,seclabel)
efivarfs on /sys/firmware/efi/efivars type efivarfs (rw,nosuid,nodev,noexec,relatime)
bpf on /sys/fs/bpf type bpf (rw,nosuid,nodev,noexec,relatime,mode=700)
/dev/mapper/cs-root on / type xfs (rw,relatime,seclabel,attr2,inode64,logbufs=8,logsize=4096k)
selinuxfs on /sys/fs/selinux type selinuxfs (rw,nosuid,noexec,relatime)
systemd-1 on /proc/sys/fs/binfmt_misc type autofs (rw,relatime,fd=29,pgrp=1,timeout=60)
mqueue on /dev/mqueue type mqueue (rw,nosuid,nodev,noexec,relatime,seclabel)
hugetlbfs on /dev/hugepages type hugetlbfs (rw,relatime,seclabel,pagesize=2M)
debugfs on /sys/kernel/debug type debugfs (rw,nosuid,nodev,noexec,relatime,seclabel)
tracefs on /sys/kernel/tracing type tracefs (rw,nosuid,nodev,noexec,relatime,seclabel)
fusectl on /sys/fs/fuse/connections type fusectl (rw,nosuid,nodev,noexec,relatime)
configfs on /sys/kernel/config type configfs (rw,nosuid,nodev,noexec,relatime)
none on /run/credentials/systemd-sysctl.service type ramfs (ro,nosuid,nodev,noexec)
none on /run/credentials/systemd-tmpfiles-setup-dev.service type ramfs (ro,nosuid,nodev)
nfsd on /proc/fs/nfsd type nfsd (rw,relatime)
/dev/mapper/cs-home on /home type xfs (rw,relatime,seclabel,attr2,inode64,logbufs=8,logsize=4096k)
/dev/sda2 on /boot type xfs (rw,relatime,seclabel,attr2,inode64,logbufs=8,logsize=4096k)
/dev/sda1 on /boot/efi type vfat (rw,relatime,fmask=0077,dmask=0077,codepage=437,allow_utime=0x0)
none on /run/credentials/systemd-tmpfiles-setup.service type ramfs (ro,nosuid,nodev)
sunrpc on /var/lib/nfs/rpc_pipefs type rpc_pipefs (rw,relatime)
tmpfs on /run/user/1000 type tmpfs (rw,nosuid,nodev,relatime,seclabel,size=174688k)
/dev/sdb1 on /data type ext4 (rw,relatime,seclabel)
```

12. Run “**lsblk -o NAME, FSTYPE, SIZE, MOUNTPOINT,LABEL**“

```
lusy@localhost.localdomain ~ (0.029s)
lsblk -o NAME,FSTYPE,SIZE,MOUNTPOINT,LABEL

NAME      FSTYPE      SIZE MOUNTPOINT LABEL
sda                   64G
└─sda1    vfat       600M /boot/efi
└─sda2    xfs        1G  /boot
└─sda3   LVM2_member 62.4G
  ├─cs-root xfs      40.6G /
  ├─cs-swap swap      2G  [SWAP]
  └─cs-home xfs      19.8G /home
sdb                   4G
└─sdb1    ext4       4G  /data      data
sr0                  1024M
```

```

lusy@localhost.localdomain ~ (0.023s)
mount | grep '^/dev'
/dev/mapper/cs-root on / type xfs (rw,relatime,seclabel,attr2,inode64,logbufs=8,logbsize=32k,
/dev/mapper/cs-home on /home type xfs (rw,relatime,seclabel,attr2,inode64,logbufs=8,logbsize=
/dev/sda2 on /boot type xfs (rw,relatime,seclabel,attr2,inode64,logbufs=8,logbsize=32k,noquot
/dev/sda1 on /boot/efi type vfat (rw,relatime,fmask=0077,dmask=0077,codepage=437,iocharset=as
/dev/sdb1 on /data type ext4 (rw,relatime,seclabel)

```

11. To automatically mount the file system at the time of boot we need the below entry to be added to the /etc/stab file

LABEL=data /data ext4 defaults 1 2

```

lusy@localhost.localdomain ~ (54.618s)
sudo vi /etc/fstab

```

```

#
# /etc/fstab
# Created by anaconda on Fri Aug  4 05:32:46 2023
#
# Accessible filesystems, by reference, are maintained under '/dev/disk/'.
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info.
#
# After editing this file, run 'systemctl daemon-reload' to update systemd
# units generated from this file.
#
/dev/mapper/cs-root      /          xfs    defaults      0 0
UUID=6a810383-8037-4807-b8bb-9a7349dbdf86 /boot      xfs    defaults      0 0
UUID=D776-2717           /boot/efi    vfat   umask=0077,shortname=winnt 0 2
/dev/mapper/cs-home      /home      xfs    defaults      0 0
/dev/mapper/cs-swap      none       swap   defaults      0 0
LABEL=data     /data      ext4    defaults      1 2

```

```

lusy@localhost.localdomain ~ (0.026s)
lsblk

```

NAME	MAJ:MIN	RM	SIZE	RO	TYPE	MOUNTPOINTS
sda	8:0	0	64G	0	disk	
└─sda1	8:1	0	600M	0	part	/boot/efi
└─sda2	8:2	0	1G	0	part	/boot
└─sda3	8:3	0	62.4G	0	part	
└─cs-root	253:0	0	40.6G	0	lvm	/
└─cs-swap	253:1	0	2G	0	lvm	[SWAP]
└─cs-home	253:2	0	19.8G	0	lvm	/home
sdb	8:16	0	4G	0	disk	
└─sdb1	8:17	0	4G	0	part	/data
sr0	11:0	1	1024M	0	rom	

Assignment 4 - Expanding a filesystem on the fly using Logical Volume Management.

Instructions

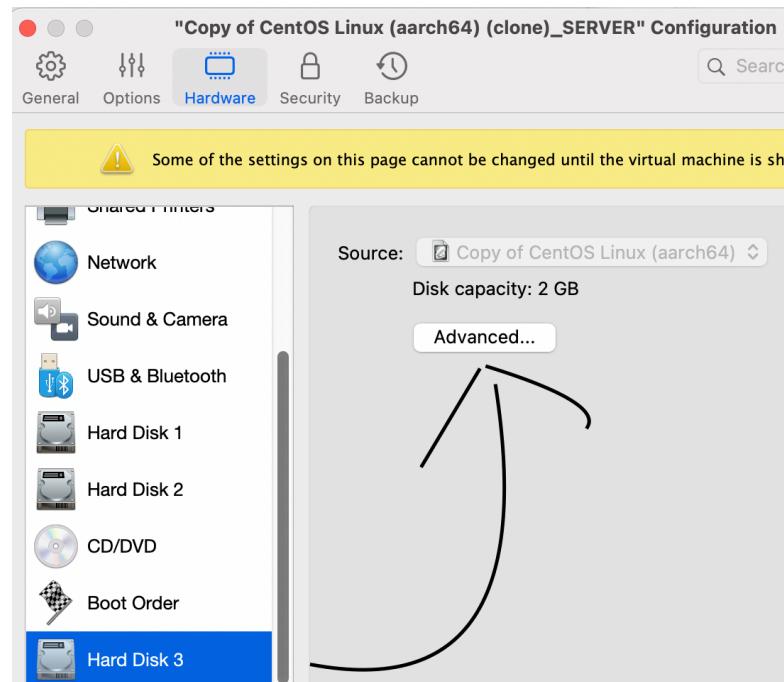
Note: install LVM using sudo yum install lvm2

Adding a new logical volume

```
lusy@localhost.localdomain ~ (5.316s)
sudo yum install lvm2

[sudo] password for lusy:
Last metadata expiration check: 0:11:25 ago on Sat 19 Aug 2023 11:27:08 PM IST.
Package lvm2-9:2.03.21-3.el9.aarch64 is already installed.
Dependencies resolved.
Nothing to do.
Complete!
```

1.1 Install a new hard drive.



1.2 Use **lsblk** command to identify the new hard drive.

```
~ (0.437s)
ssh lusy@10.211.55.6

lusy@localhost.localdomain ~ (0.034s)
lsblk

| NAME      | MAJ:MIN | RM | SIZE  | R0 | TYPE | MOUNTPOINTS |
|-----------|---------|----|-------|----|------|-------------|
| sda       | 8:0     | 0  | 64G   | 0  | disk |             |
| └─sda1    | 8:1     | 0  | 600M  | 0  | part | /boot/efi   |
| └─sda2    | 8:2     | 0  | 1G    | 0  | part | /boot       |
| └─sda3    | 8:3     | 0  | 62.4G | 0  | part |             |
| └─cs-root | 253:0   | 0  | 40.6G | 0  | lvm  | /           |
| └─cs-swap | 253:1   | 0  | 2G    | 0  | lvm  | [SWAP]      |
| └─cs-home | 253:2   | 0  | 19.8G | 0  | lvm  | /home       |
| sdb       | 8:16    | 0  | 4G    | 0  | disk |             |
| └─sdb1    | 8:17    | 0  | 4G    | 0  | part | /data       |
| sdc       | 8:32    | 0  | 2G    | 0  | disk |             |
| sr0       | 11:0    | 1  | 1024M | 0  | rom  |             |


```

1.3 Create Physical Volume from the hard drive using "**pvcreate /dev/sdc**". Assuming **/dev/sde** is the new hard drive created.

```
lusy@localhost.localdomain ~ (0.18s)
sudo pvcreate /dev/sdc

Physical volume "/dev/sdc" successfully created.
```

1.4 You can list all the PV with the command: **pvscan**.

```
lusy@localhost.localdomain ~ (0.142s)
sudo pvscan

PV /dev/sda3   VG cs                  lvm2 [62.41 GiB / 0    free]
PV /dev/sdc          lvm2 [2.00 GiB]
Total: 2 [64.41 GiB] / in use: 1 [62.41 GiB] / in no VG: 1 [2.00 GiB]
```

1.5 If you want to display more information about any specific PV, let's say **/dev/sd**c run “**pvdisplay /dev/sdc**”

```
lousy@localhost.localdomain ~ (0.144s)
sudo pvdisplay /dev/sdc
--- Physical volume ---
PV Name              /dev/sdc
VG Name              share
PV Size              2.00 GiB / not usable 4.00 MiB
Allocatable          yes
PE Size              4.00 MiB
Total PE             511
Free PE              230
Allocated PE         281
PV UUID              nGfLYr-7dAY-i27U-BWZ3-cHJ0-eF6e-6PB0wr
```

1.6 Create Volume Group named "share" using **sudo vgcreate share / dev/sdc**.

```
lousy@localhost.localdomain ~ (0.182s)
sudo vgcreate share /dev/sdc
Volume group "share" successfully created
```

1.7 you can list all the VGs with the command: **vgscan**

```
lousy@localhost.localdomain ~ (0.166s)
sudo vgscan
Found volume group "share" using metadata type lvm2
Found volume group "cs" using metadata type lvm2
```

1.8 You can display more information about any specific VG, such as share with the command: **vgdisplay share**.

```
lousy@localhost.localdomain ~ (0.178s)
sudo vgdisplay share
--- Volume group ---
VG Name          share
System ID
Format           lvm2
Metadata Areas   1
Metadata Sequence No 1
VG Access        read/write
VG Status         resizable
MAX LV            0
Cur LV            0
Open LV           0
Max PV            0
Cur PV            1
Act PV            1
VG Size           <2.00 GiB
PE Size           4.00 MiB
Total PE          511
Alloc PE / Size   0 / 0
Free PE / Size    511 / <2.00 GiB
VG UUID          s7ss9d-mi9g-JPXs-XX14-t2oH-qb0X-p7uGmQ
```

1.9 Create Logical Volume: To create a 100MB LV mylv from VG share use the command "**sudo lvcreate --size 100M --name mylv share**"

```
lousy@localhost.localdomain ~ (0.213s)
sudo lvcreate --size 100M --name mylv share
Logical volume "mylv" created.
```

1.10 create another LV mylv1 of size 1GB from VG share with the command: " **sudo lvcreate --size 1G --name mylv1 share**"

```
lusy@localhost.localdomain ~ (0.154s)
sudo lvcreate --size 1G --name mylv1 share
Logical volume "mylv1" created.
```

1.11 list all the LVs with the commands: **lvscan** or **lvs**.

```
lusy@localhost.localdomain ~ (0.165s)
sudo lvscan

ACTIVE          '/dev/share/mylv' [100.00 MiB] inherit
ACTIVE          '/dev/share/mylv1' [1.00 GiB] inherit
ACTIVE          '/dev/cs/swap' [2.04 GiB] inherit
ACTIVE          '/dev/cs/home' [19.80 GiB] inherit
ACTIVE          '/dev/cs/root' [40.56 GiB] inherit
```

1.12 You can also display more information about any specific LV with the command: " **sudo lvdisplay VG_NAME/LV_NAME**".

1.13 The LVs are available as /dev/VG NAME/LV NAME. You should be able to see the 2 Vs created so far.

```
lusy@localhost.localdomain ~ (0.186s)
sudo lvdisplay /dev/share

--- Logical volume ---
LV Path          /dev/share/mylv
LV Name          mylv
VG Name          share
LV UUID          YIgZ5G-JKuR-eEQq-noXi-YEGX-cRkj-B7ZMEd
LV Write Access  read/write
LV Creation host, time localhost.localdomain, 2023-08-19 23:45:50 +0530
LV Status        available
# open           1
LV Size          100.00 MiB
Current LE      25
Segments         1
Allocation       inherit
Read ahead sectors auto
- currently set to 8192
Block device    253:3
```

```
--- Logical volume ---
LV Path          /dev/share/mylv1
LV Name          mylv1
VG Name          share
LV UUID          VRz6ED-9GQx-tW0g-bsbx-Pj6L-8MiR-WG830n
LV Write Access  read/write
LV Creation host, time localhost.localdomain, 2023-08-19 23:46:01 +0530
LV Status        available
# open           1
LV Size          1.00 GiB
Current LE       256
Segments         1
Allocation       inherit
Read ahead sectors auto
- currently set to 8192
Block device    253:4
```

1.14 Format the Logical Volumes

mkfs-text4 /dev/share/mylv

```
lusy@localhost.localdomain ~ (0.079s)
sudo mkfs.ext4 /dev/share/mylv

mke2fs 1.46.5 (30-Dec-2021)
Discarding device blocks: done
Creating filesystem with 102400 1k blocks and 25584 inodes
Filesystem UUID: 94424546-8000-46ef-8479-bc341640d663
Superblock backups stored on blocks:
      8193, 24577, 40961, 57345, 73729

Allocating group tables: done
Writing inode tables: done
Creating journal (4096 blocks): done
Writing superblocks and filesystem accounting information: done
```

mkfs-text4 /dev/share/mylv1

```
lusy@localhost.localdomain ~ (0.078s)
sudo mkfs.ext4 /dev/share/mylv1

mke2fs 1.46.5 (30-Dec-2021)
Discarding device blocks: done
Creating filesystem with 262144 4k blocks and 65536 inodes
Filesystem UUID: aaf93c14-e998-4890-a994-0449712576e9
Superblock backups stored on blocks:
      32768, 98304, 163840, 229376

Allocating group tables: done
Writing inode tables: done
Creating journal (8192 blocks): done
Writing superblocks and filesystem accounting information: done
```

1.15 create mount points where you want to mount these LVs

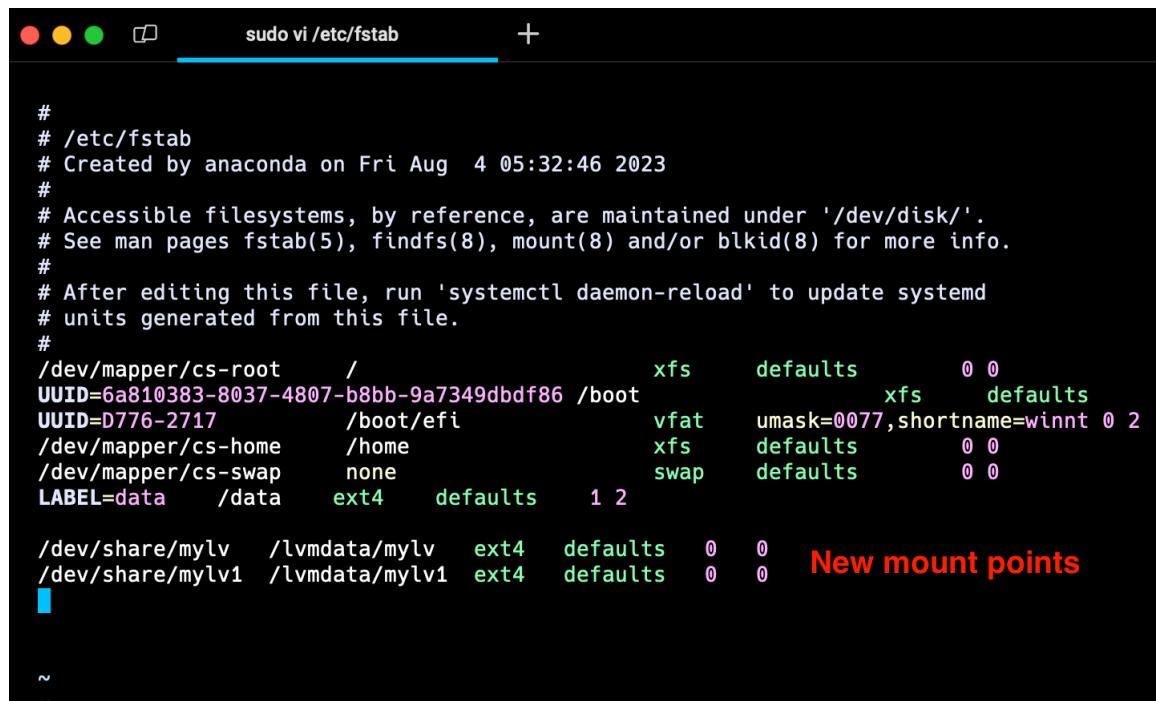
1.16 add an appropriate entry to the **/etc/fstab** file, and mount the filesystem.

sudo mkdir -pv /lvmda

```
lusy@localhost.localdomain ~ (0.045s)
sudo mkdir -pv /lvmda/mylv
mkdir: created directory '/lvmda'
mkdir: created directory '/lvmda/mylv'
```

```
lusy@localhost.localdomain ~ (0.059s)
sudo mkdir -pv /lvmda/mylv1
mkdir: created directory '/lvmda/mylv1'
```

```
lusy@localhost.localdomain ~ (51.98s)
sudo vi /etc/fstab
```



```
# /etc/fstab
# Created by anaconda on Fri Aug  4 05:32:46 2023
#
# Accessible filesystems, by reference, are maintained under '/dev/disk/'.
# See man pages fstab(5), findfs(8), mount(8) and/or blkid(8) for more info.
#
# After editing this file, run 'systemctl daemon-reload' to update systemd
# units generated from this file.
#
/dev/mapper/cs-root      /          xfs    defaults      0  0
UUID=6a810383-8037-4807-b8bb-9a7349dbdf86 /boot      xfs    defaults      0  0
UUID=D776-2717            /boot/efi    vfat   umask=0077,shortname=winnt 0  2
/dev/mapper/cs-home       /home       xfs    defaults      0  0
/dev/mapper/cs-swap       none        swap   defaults      0  0
LABEL=data     /data      ext4    defaults    1  2
/dev/share/mylv  /lvmda/mylv  ext4    defaults    0  0  New mount points
/dev/share/mylv1 /lvmda/mylv1 ext4    defaults    0  0
```

```
lusy@localhost.localdomain ~ (0.065s)
sudo mount -a
mount: (hint) your fstab has been modified, but systemd still uses
      the old version; use 'systemctl daemon-reload' to reload.
```

```
lusy@localhost.localdomain ~ (7.6s)
systemctl daemon-reload
===== AUTHENTICATING FOR org.freedesktop.systemd1.reload-daemon =====
Authentication is required to reload the systemd state.
Authenticating as: root
Password:
===== AUTHENTICATION COMPLETE =====
```

```
lusy@localhost.localdomain ~ (0.03s)
lsblk

| NAME          | MAJ:MIN | RM | SIZE  | R0 | TYPE | MOUNTPOINTS  |
|---------------|---------|----|-------|----|------|--------------|
| sda           | 8:0     | 0  | 64G   | 0  | disk |              |
| └─sda1        | 8:1     | 0  | 600M  | 0  | part | /boot/efi    |
| └─sda2        | 8:2     | 0  | 1G    | 0  | part | /boot        |
| └─sda3        | 8:3     | 0  | 62.4G | 0  | part |              |
| └─cs-root     | 253:0   | 0  | 40.6G | 0  | lvm  | /            |
| └─cs-swap     | 253:1   | 0  | 2G    | 0  | lvm  | [SWAP]       |
| └─cs-home     | 253:2   | 0  | 19.8G | 0  | lvm  | /home        |
| sdb           | 8:16    | 0  | 4G    | 0  | disk |              |
| └─sdb1        | 8:17    | 0  | 4G    | 0  | part | /data        |
| sdc           | 8:32    | 0  | 2G    | 0  | disk |              |
| └─share-mylv  | 253:3   | 0  | 100M  | 0  | lvm  | /lvmda       |
| └─share-mylv1 | 253:4   | 0  | 1G    | 0  | lvm  | /lvmda/mylv1 |
| sr0           | 11:0    | 1  | 1024M | 0  | rom  |              |


```